Homework 2

CS420 Machine learning 2020 Spring*
Department of Computer Science and Engineering
Shanghai Jiao Tong University

Submission deadline: 20:00, May 16, 2020, Saturday

Submission to:

Please submit your homework in pdf/doc format to Canvas platform.

1 (10 points) PCA algorithm

Give at least two algorithms that could take data set $X = \{x_1, \ldots, x_N\}$, $x_t \in \mathbb{R}^{n\times 1}, \forall t$ as input, and output the first principal component $\mathbf{w}$. Specify the computational details of the algorithms, and discuss the advantages or limitations of the algorithms.

2 (10 points) Factor Analysis (FA)

Calculate the Bayesian posterior $p(\mathbf{y}|\mathbf{x})$ of the Factor Analysis model $\mathbf{x} = \mathbf{A}\mathbf{y} + \mu + \mathbf{e}$, with $p(\mathbf{x}|\mathbf{y}) = G(\mathbf{x}|\mathbf{A}\mathbf{y} + \mu, \Sigma_e)$, $p(\mathbf{y}) = G(\mathbf{y}|0, \Sigma_y)$, where $G(\mathbf{z}|\mu, \Sigma)$ denotes Gaussian distribution density with mean μ and covariance matrix Σ.

3 (10 points) Independent Component Analysis (ICA)

Explain why maximizing non-Gaussianity could be used as a principle for ICA estimation.

4 (50 points) Dimensionality Reduction by FA

Consider the following Factor Analysis (FA) model,

$$\mathbf{x} = \mathbf{A}\mathbf{y} + \mu + \mathbf{e}, \tag{1}$$

$$p(\mathbf{x}|\mathbf{y}) = G(\mathbf{x}|\mathbf{A}\mathbf{y} + \mu, \sigma^2\mathbf{I}), \tag{2}$$

$$p(\mathbf{y}) = G(\mathbf{y}|0, \mathbf{I}), \tag{3}$$

where the observed variable $\mathbf{x} \in \mathcal{R}^n$, the latent variable $\mathbf{y} \in \mathcal{R}^m$, and $G(\mathbf{z}|\mu, \Sigma)$ denotes Gaussian distribution density with mean μ and covariance matrix Σ. Write a report on experimental comparisons on model selection performance by BIC, AIC on selecting the number of latent factors, i.e., $\mathbf{dim}(\mathbf{y}) = m$.

Specifically, you need to randomly generate datasets based on FA, by varying some setting values, e.g., sample size N, dimensionality n and m, noise level σ^2, and so on. For example, set $N = 100, n = 10, m = 3, \sigma^2 = 0.1, \mu = 0$, and assign values for $\mathbf{A} \in \mathcal{R}^{n\times m}$. The generation process is as follows:

(1) Randomly sample a $\mathbf{y}_t$ from Gaussian density $G(\mathbf{y}|0, \mathbf{I})$, with $\mathbf{dim}(\mathbf{y}) = m = 3$;

*tushikui@sjtu.edu.cn

(2) Randomly sample a noise vector $\mathbf{e}_t$ from Gaussian density $G(\mathbf{e}|0, \sigma^2\mathbf{I})$, with $\sigma^2 = 0.1$, $\mathbf{e}_t \in \mathcal{R}^n$;

(3) Get $\mathbf{x}_t = \mathbf{A}\mathbf{y}_t + \mu + \mathbf{e}_t$.

Collect all the $\mathbf{x}_t$ as the dataset $X = \{\mathbf{x}_t\}_{t=1}^N$.

The two-stage model selection process for BIC, AIC is as follows:

Stage 1: Run EM algorithm on each dataset X for $m = 1, ..., M$, and calculate the log-likelihood value $\ln[p(X|\hat{\Theta}_m)]$, where $\hat{\Theta}_m$ is the maximum likelihood estimate for parameters;

Stage 2: Select the optimal m^* by

$$m^* = \arg\max_{m=1,...,M} J(m), \tag{4}$$

$$J_{AIC}(m) = \ln[p(X|\hat{\Theta}_k)] - d_m \tag{5}$$

$$J_{BIC}(m) = \ln[p(X|\hat{\Theta}_k)] - \frac{\ln N}{2} d_m \tag{6}$$

You may set $M = 5$, if you generate the dataset X based on $n = 10, m = 3$.

The following codes might be useful.

Python: `https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.FactorAnalysis.html#sklearn.decomposition.FactorAnalysis`

5 (20 points) Spectral clustering

Use experiments to demonstrate that when spectral clustering works well, when it would fail. Summarize your results.

The following codes might be helpful.
Python: `https://scikit-learn.org/stable/modules/generated/sklearn.cluster.SpectralClustering.html`